

[Open in app](#)

Tim van Baarsen

[Following](#)

171 Followers

[About](#)

Programmatically create Kafka topics using Spring Kafka

**Tim van Baarsen** Apr 10, 2020 · 5 min read

Since the introduction of the `AdminClient` in the Kafka Clients library (version 0.11.0.0), we can create topics programmatically. Spring Kafka is leveraging the Kafka `AdminClient` to create Kafka topics programmatically even easier! Learn both about how to use it but also how to avoid some pitfalls 😊.



Spring Boot will autoconfigure a `AdminClient` Spring Bean in your application context which will automatically add topics for all beans of type `NewTopic`.

Spring Kafka version 2.3 introduced a `TopicBuilder` class to make the creation of such beans even more convenient!

[Open in app](#)

used by Spring as a source of bean definitions.

All methods annotated with `@Bean` : will return an object that should be registered as a bean in the Spring application context. Spring Kafka will automatically add topics for all beans of type `NewTopic`

```
1 package io.stockgeeks.kafka.config;
2
3 import org.apache.kafka.clients.admin.NewTopic;
4 import org.springframework.context.annotation.Bean;
5 import org.springframework.context.annotation.Configuration;
6 import org.springframework.kafka.config.TopicBuilder;
7
8 @Configuration
9 public class KafkaTopicConfiguration {
10
11     @Bean
12     public NewTopic topicExample() {
13         return TopicBuilder.name("my-first-kafka-topic")
14             .partitions(6)
15             .replicas(3)
16             .build();
17     }
18 }
```

KafkaTopicConfiguration.java hosted with by GitHub

[view raw](#)

Next to the name of the Kafka topic name you can specify:

- the number of partitions for the topic
- the number of replicas for the topic
- assign replicas

To create a **compact** Kafka topic:

```
1 package io.stockgeeks.kafka.config;
2
3 import org.apache.kafka.clients.admin.NewTopic;
4 import org.springframework.context.annotation.Bean;
5 import org.springframework.context.annotation.Configuration;
6 import org.springframework.kafka.config.TopicBuilder;
```

[Open in app](#)

```
10
11     @Bean
12     public NewTopic compactTopicExample() {
13         return TopicBuilder.name("my-first-compact-kafka-topic")
14             .partitions(6)
15             .replicas(3)
16             .compact()
17             .build();
18     }
19 }
```

KafkaTopicConfiguration.java hosted with by GitHub

[View raw](#)

Example of passing specific Kafka **topic configuration properties** (in this case to configure compression:

```
1  import org.apache.kafka.clients.admin.NewTopic;
2  import org.apache.kafka.common.config.TopicConfig;
3  import org.springframework.context.annotation.Bean;
4  import org.springframework.context.annotation.Configuration;
5  import org.springframework.kafka.config.TopicBuilder;
6
7  @Configuration
8  public class KafkaTopicConfiguration {
9
10     @Bean
11     public NewTopic topicWithCompressionExample() {
12         return TopicBuilder.name("kafka-topic-with-compression")
13             .partitions(6)
14             .replicas(3)
15             .config(TopicConfig.COMPRESSION_TYPE_CONFIG, "zstd")
16             .build();
17     }
18 }
```

KafkaTopicConfiguration.java hosted with by GitHub

[View raw](#)

Caveats and pitfalls

There are some caveats and pitfalls to take into account when creating topics programmatically. Let's go over them one by one:

- Not using Spring Boot?

[Open in app](#)

The default number of partitions and replication factor in the topic are:

- Replication factor in a single node Kafka Cluster for local development

Not using Spring Boot?

In case you are not using Spring Boot, you have to configure the `KafkaAdmin` bean yourself to automatically add topics for all beans of type `NewTopic`

```
1  package io.stockgeeks.kafka.config;
2
3  import org.apache.kafka.clients.admin.AdminClientConfig;
4  import org.springframework.beans.factory.annotation.Value;
5  import org.springframework.context.annotation.Bean;
6  import org.springframework.context.annotation.Configuration;
7  import org.springframework.kafka.core.KafkaAdmin;
8
9  import java.util.HashMap;
10 import java.util.Map;
11
12 @Configuration
13 public class KafkaTopicConfiguration {
14
15     @Value(value = "${kafka.bootstrapServers:localhost:9092}")
16     private String bootstrapServers;
17
18     @Bean
19     public KafkaAdmin kafkaAdmin() {
20         Map<String, Object> configs = new HashMap<>();
21         // Depending on you Kafka Cluster setup you need to configure
22         // additional properties!
23         configs.put(AdminClientConfig.BOOTSTRAP_SERVERS_CONFIG, bootstrapServers);
24         return new KafkaAdmin(configs);
25     }
26 }
```

Automatically increases the number of partitions

From the Spring Kafka documentation:

If the broker supports it (1.0.0 or higher), the admin increases the number of partitions if it is found that an existing topic has fewer partitions than the `NewTopic.numPartitions`.

[Open in app](#)

- name “**already-existing-kafka-topic**”
- with three partitions

Show details about the topic using the command: `kafka-topics.sh` (Part of the Apache Kafka distribution)

```
bin/kafka-topics.sh --bootstrap-server localhost:9092 --zookeeper
zookeeper:2181 --topic already-existing-kafka-topic --describe
```

Output showing us the number of partitions is 3:

```
Topic: already-existing-kafka-topic      PartitionCount: 3
ReplicationFactor: 1    Configs:
    Topic: already-existing-kafka-topic      Partition: 0
Leader: 1001    Replicas: 1001    Isr: 1001
    Topic: already-existing-kafka-topic      Partition: 1
Leader: 1001    Replicas: 1001    Isr: 1001
    Topic: already-existing-kafka-topic      Partition: 2
Leader: 1001    Replicas: 1001    Isr: 1001
```

Now you configured six partitions for the topic using the `TopicBuilder` and start your application:

Although the topic already exists, the number of partitions of the topic is increased to six!

```
bin/kafka-topics.sh --bootstrap-server localhost:9092 --zookeeper
zookeeper:2181 --topic already-existing-kafka-topic --describe
```

Output showing the number of partitions increased:

```
Topic: already-existing-kafka-topic      PartitionCount: 6
ReplicationFactor: 1    Configs:
    Topic: already-existing-kafka-topic      Partition: 0
```

[Open in app](#)

```

Topic: already-existing-kafka-topic Partition: 2
Leader: 1001 Replicas: 1001 Isr: 1001
Topic: already-existing-kafka-topic Partition: 3
Leader: 1001 Replicas: 1001 Isr: 1001
Topic: already-existing-kafka-topic Partition: 4
Leader: 1001 Replicas: 1001 Isr: 1001
Topic: already-existing-kafka-topic Partition: 5
Leader: 1001 Replicas: 1001 Isr: 1001

```

Be aware:

- If partitions are increased for a topic, and the producer is using a key to produce messages, the partition logic or ordering of the messages will be affected!
- The number of partitions for a Kafka topic can only be increased.

The default number of partitions and replication count in the TopicBuilder

By default, the values for both the **partitions** and **replicas** in the TopicBuilder are one! For a local development environment, this is a sensible default (because, in most cases, you just run a single node Kafka cluster). But production Kafka applications need topics with a proper number of partitions and replicas to be able to balance the load, scale, and be fault-tolerant.

There is no silver bullet for the number of partitions of your Kafka topic. It depends on your use case. On the Confluent blog, you can find a good read about [how to choose the number of topic partitions](#).

If you would like to know more about the basics of partitions and replicas read my previous blog post: “[Head First Kafka: The basics of producing data to Kafka explained using a conversation](#).”

Replication factor in a single node Kafka cluster for local development

From the Kafka documentation: The partitions of the log are distributed over the servers in the Kafka cluster with each server handling data and requests for a share of the partitions. Each partition is replicated across a configurable number of servers for fault tolerance.

The **replication factor** defines the number of copies for each message produced to a Kafka topic. The replication factor can be defined at the topic level (like we do here in the Java config). Replicas are distributed evenly among Kafka brokers in a cluster.

[Open in app](#)

you will run into this error:

```
2020-04-10 10:27:59.925 ERROR 3013 --- [           main]
o.springframework.kafka.core.KafkaAdmin : Could not configure
topics

org.springframework.kafka.KafkaException: Failed to create topics;
nested exception is
org.apache.kafka.common.errors.InvalidReplicationFactorException:
Replication factor: 3 larger than available brokers: 1.
```

Long story short: when you run a single node Kafka cluster for local development, there are no other brokers to replicate the data. So you can only configure a replication factor of one!

Since you don't want to configure a replication factor of one for non-development environments, you need to make the number of replicas configurable!

Takeaways from my blogpost

- Programmatically creating Kafka topics is powerful but be aware of the pitfalls.
- Plan the number of partitions and replicas for your topic ahead based on your use-case.
- For a local development single node Kafka cluster, you can only configure a replication factor of one.
- The use of the AdminClient might be restricted on your Kafka cluster that makes it not possible to programmatically creating Kafka topics

Keep on learning

Use useful links related to this blogpost:

- [Javadoc Kafka AdminClient](#)
- [Spring Kafka Documentation about configuring topics](#)
- [Anatomy of a topic](#)

[Open in app](#)

Tap the  button if you found this article useful!

Any questions or feedback? Reach out to me on Twitter: [@TimvanBaarsen](#)

[Kafka](#) [Spring](#) [Spring Boot](#) [Spring Kafka](#) [Java](#)

[About](#) [Help](#) [Legal](#)

Get the Medium app

