

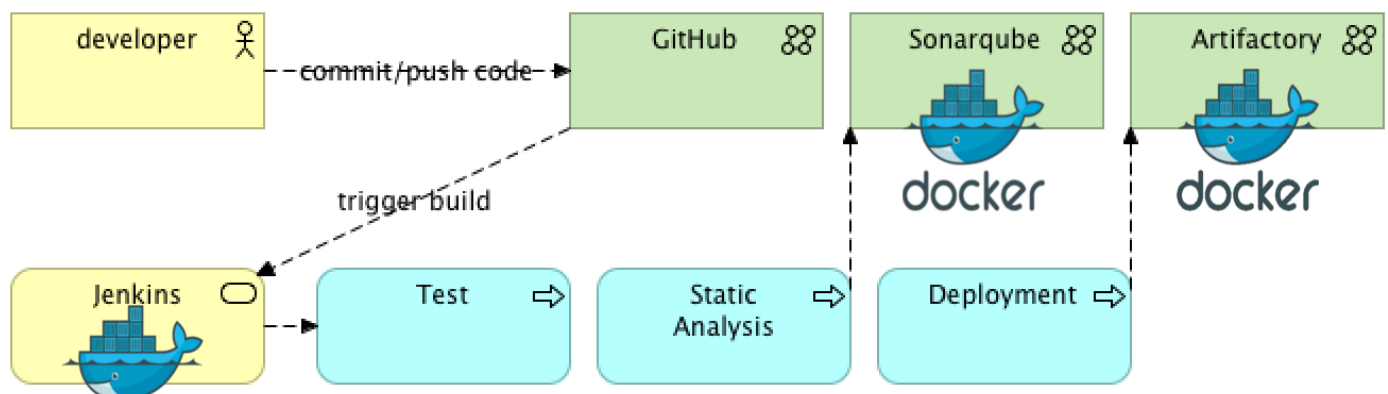


Dockerizing Jenkins, Part 2: Deployment With Maven and JFrog Artifactory

by **Kayan Azimov** · Jul. 31, 17 · DevOps Zone

Download the blueprint that can take a company of any maturity level all the way up to enterprise-scale continuous delivery using a combination of Atomic Release Automation, Atomic's 20+ years of business automation experience, and the proven tools and practices the company is already leveraging.

In the first part of this tutorial, we looked at how to dockerize installation of the Jenkins plugins, Java and Maven tool setup in Jenkins 2, and created a declarative build pipeline for a Maven project with test and SonarQube stages. In this part, we will focus on deployment.



Couldn't we simply add another stage for deployment in part 1, you may ask? Well, in fact, deployment requires quite a few steps to be taken, including Maven pom and settings file configuration, artifact repository availability, repository credentials encryption, etc. Let's add them to the list and then implement them step by step like we did in the previous session.

- Running JFrog Artifactory on Docker.
- Configuring the Maven pom file.
- Configuring the Maven settings file.
- Using Config File Provider Plugin for persistence of Maven settings.
- Dockerizing the installation and configuration process.

If you are already familiar with first part of this tutorial, created your project from the scratch and using your own repository, then you can just follow the steps as we go further, otherwise, if you are starting now, you can just clone/fork the work we did in the last example and then add changes as they follow in the tutorial:

```
git clone https://github.com/kenych/jenkins-docker-pipeline_tutorial11 && cd jenkins-docker
```

Please note all steps have been tested on MacOS Sierra and Docker version 17.05.0-ce and you should change them accordingly if you are using MS-DOS, FreeBSD, etc.

The script above is going to take a while as it is downloading Java 7, Java 8, Maven, SonarQube, and Jenkins Docker images, so please be patient. Once done, you should have Jenkins and Sonar up and running as we created in part 1:

The screenshot shows the Jenkins web interface at `localhost:8080/job/berlin-clock/`. The interface includes a sidebar with navigation links: [Back to Dashboard](#), [Status](#), [Changes](#), [Build Now](#), [Delete Pipeline](#), [Configure](#), [Full Stage View](#), and [Pipeline Syntax](#). The main content area displays the 'Pipeline View' for the 'berlin-clock' pipeline. It shows a 'Stage View' with a table of stage times:

Stage	Declarative: Checkout SCM	Declarative: Tool Install	install and sonar parallel
Average stage times:	4s	5s	54s
#1	4s	5s	54s

Below the stage view, there is a 'Permalinks' section with a link: [Last build \(#1\), 1 min 7 sec ago](#). On the left, there is a 'Build History' section showing a search bar and a list of builds, including build #1 from Jul 16, 2017 6:24 PM.

The screenshot shows the JFrog Artifactory web interface at `localhost:9000/projects`. The interface includes a sidebar with navigation links: [Projects](#), [Issues](#), [Rules](#), [Quality Profiles](#), and [Quality Gates](#). The main content area displays the 'Quality Gate' for the 'berlin-clock' project. It shows a table of quality metrics:

Metric	Value
Reliability	A
Security	A
Maintainability	A
Coverage	0.0%
Duplications	0.0%
Java	XS 301

The 'Quality Gate' is shown as 'Passed'.

If you got errors about some port being busy, just use the free ports from your host, as I explain [here](#). Otherwise, you can use dynamic ports.

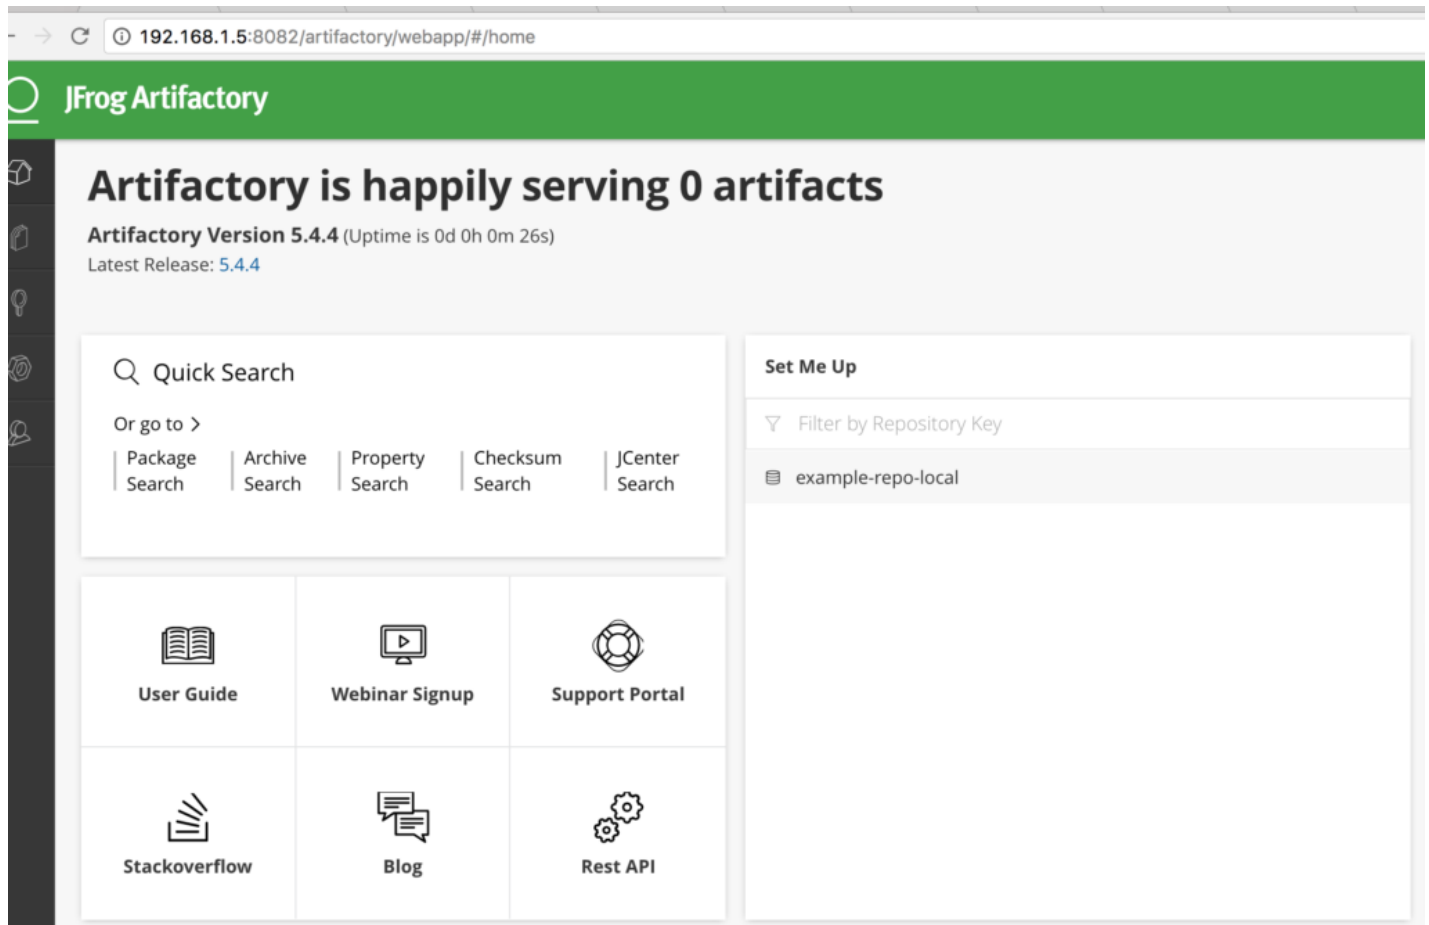
Chapter 1: Running JFrog Artifactory on Docker

So let's look at the first step. Obviously, if we want to test the deployment in our example, we need some place to deploy our artifacts to. We are going to use a limited open source version of JFrog

Artifactory called “artifactory oss.” Let’s run it on Docker to see how easy it is to have your own artifact repo. The port 8081 on my machine was busy, so I had to run it on 8082. You should do so according to free ports available on your machine:

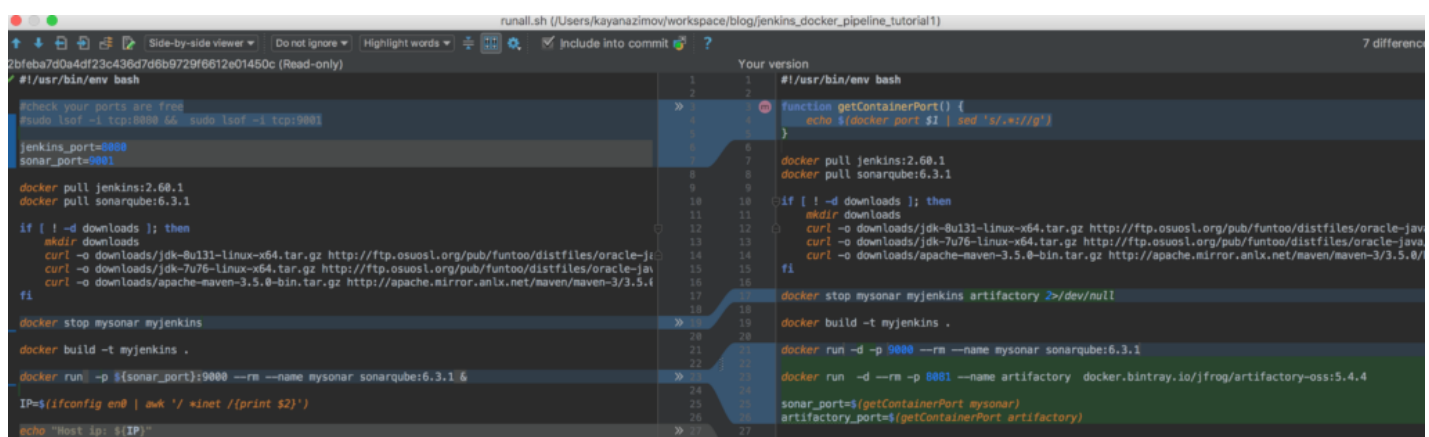
```
docker run --rm -p 8082:8081 --name artifactory docker.bintray.io/jfrog/artifactory-oss:5.
```

Normally you would run Artifactory with volumes to preserve its state (the jar files in our case), but for the sake of our tutorial which is about Dockerizing Jenkins, we will simply run it on “in memory” fashion. Once it is up and running, we should be ready to deploy our artifacts to it.



Please pay attention to the name of the path for the default repository it has created; it was example-repo-local in my case, and we will refer to it very soon.

By the time I finished part two, I improved the runall script; now it uses dynamic ports and should run with no issues if some of the ports used by SonarQube, Jenkins or Artifactory were busy (8080, 8081, 9000, etc):



```

20 20 if [ ! -d m2deps ]; then
21 21     mkdir m2deps
22 22 fi
23 23
24 24 docker run -p ${jenkins_port}:8080 -v `pwd`/downloads:/var/jenkins_home/downloads \
25 25     -v `pwd`/jobs:/var/jenkins_home/jobs/ \
26 26     -v `pwd`/m2deps:/var/jenkins_home/.m2/repository/ --rm --name myjenkins \
27 27     -e SONARQUBE_HOST=http://${IP}:${sonar_port} \
28 28     -e ARTIFACTORY_URL=http://${IP}:${artifactory_port}/artifactory/example-repo-local \
29 29     myjenkins:latest
30 30
31 31 echo "Sonarqube is running at ${IP}:${sonar_port}"
32 32 echo "Artifactory is running at ${IP}:${artifactory_port}"
33 33 echo "Jenkins is running at ${IP}:${getContainerPort myjenkins}"
34 34
35 35
36 36
37 37
38 38
39 39
40 40
41 41
42 42
43 43
44 44
45 45
46 46
47 47
48 48
49 49
50 50
51 51
52 52
53 53
54 54
55 55
56 56
57 57
58 58
59 59
60 60
61 61
62 62
63 63
64 64
65 65
66 66
67 67
68 68
69 69
70 70
71 71
72 72
73 73
74 74
75 75
76 76
77 77
78 78
79 79
80 80
81 81
82 82
83 83
84 84
85 85
86 86
87 87
88 88
89 89
90 90
91 91
92 92
93 93
94 94
95 95
96 96
97 97
98 98
99 99
100 100

```

So let's switch to updated script. It will just stop all containers, build a Jenkins image, and then rerun with dynamic ports:

```

1  #!/usr/bin/env bash
2
3  function getContainerPort() {
4      echo $(docker port $1 | sed 's/.*://g')
5  }
6
7  docker pull jenkins:2.60.1
8  docker pull sonarqube:6.3.1
9
10 if [ ! -d downloads ]; then
11     mkdir downloads
12     curl -o downloads/jdk-8u131-linux-x64.tar.gz http://ftp.osuosl.org/pub/funtoo/distfiles/
13     curl -o downloads/jdk-7u76-linux-x64.tar.gz http://ftp.osuosl.org/pub/funtoo/distfiles/
14     curl -o downloads/apache-maven-3.5.0-bin.tar.gz http://apache.mirror.anlx.net/maven/m
15 fi
16
17 docker stop mysonar myjenkins artifactory 2>/dev/null
18
19 docker build -t myjenkins .
20
21 docker run -d -p 9000 --rm --name mysonar sonarqube:6.3.1
22
23 docker run -d --rm -p 8081 --name artifactory docker.bintray.io/jfrog/artifactory-oss:5.
24
25 sonar_port=$(getContainerPort mysonar)
26 artifactory_port=$(getContainerPort artifactory)
27
28 IP=$(ifconfig en0 | awk '/ *inet /{print $2}')
29
30 if [ ! -d m2deps ]; then
31     mkdir m2deps
32 fi
33
34 docker run -d -p 8080 -v `pwd`/downloads:/var/jenkins_home/downloads \
35     -v `pwd`/jobs:/var/jenkins_home/jobs/ \

```

```

36 -v `pwd`/m2deps:/var/jenkins_home/.m2/repository/ --rm --name myjenkins \
37 -e SONARQUBE_HOST=http://${IP}:${sonar_port} \
38 -e ARTIFACTORY_URL=http://${IP}:${artifactory_port}/artifactory/example-repo-local \
39 myjenkins:latest
40
41 echo "Sonarqube is running at ${IP}:${sonar_port}"
42 echo "Artifactory is running at ${IP}:${artifactory_port}"
43 echo "Jenkins is running at ${IP}:${getContainerPort myjenkins}"

```

Run it in debug mode to see what is happening:

```

1  bash +x runall.sh
2  2.60.1: Pulling from library/jenkins
3  Digest: sha256:fa62fceb220e7545d1791e6eea6759b4c3bdba246dd839289f2b28b653e72
4  Status: Image is up to date for jenkins:2.60.1
5  6.3.1: Pulling from library/sonarqube
6  Digest: sha256:d5f7bb8aeca46da054bf28d111e5a27f1378188b427db64cc9fb392e1a8d80a
7  Status: Image is up to date for sonarqube:6.3.1
8  mysonar
9  myjenkins
10 artifactory
11 Sending build context to Docker daemon 365.1MB
12 Step 1/6 : FROM jenkins:2.60.1
13 ---&gt; f426a52bafa9
14 Step 2/6 : MAINTAINER Kayan Azimov
15 ---&gt; Using cache
16 ---&gt; 760e7bb0f335
17 Step 3/6 : ENV JAVA_OPTS "-Djenkins.install.runSetupWizard=false"
18 ---&gt; Using cache
19 ---&gt; e3dbac0834cd
20 Step 4/6 : COPY plugins.txt /usr/share/jenkins/ref/plugins.txt
21 ---&gt; Using cache
22 ---&gt; 39966bece010
23 Step 5/6 : RUN /usr/local/bin/install-plugins.sh &lt; /usr/share/jenkins/ref/plugi
24 ---&gt; 987bdeca2517
25 Step 6/6 : COPY groovy/* /usr/share/jenkins/ref/init.groovy.d/
26 ---&gt; Using cache
27 ---&gt; e5ec6b7f49aa
28 Successfully built e5ec6b7f49aa
29 Successfully tagged myjenkins:latest
30 85d8716d3c7fd6272b6915d977daa37dbe9e4ece0f5c367dd9798fbfca272b2d
31 5fc3f649d3ba5c47df48a54b761d432611ed18872aa727114cdcf1a0f30cac0c
32 21e98466cb4f6e78817d9869d39bda57f475d7174f565031f168eb836118d701
33 Sonarqube is running at 192.168.1.2:32836
34 Artifactory is running at 192.168.1.2:32837
35 Jenkins is running at 192.168.1.2:32838

```

We now see the ports for all running containers in the logs, so we can access any of them if we like. As you may have noticed, our shell script has become a bit too complicated; this is just to run three containers, so it means that next time, I should probably switch to using docker compose instead!

We obviously need to update our pipeline which we created in part 1, add a deployment step to it, push the updated Jenkins file, and build the job. If you are using your own repo, alternatively, just amend the existing job as below by using Replthe ay button on the last successful build and then run it:

```

1  pipeline {
2      agent any
3
4      tools {
5          jdk 'jdk8'
6          maven 'maven3'
7      }
8
9      stages {
10         stage('install and sonar parallel') {
11             steps {
12                 parallel(
13                     install: {
14                         sh "mvn -U clean test cobertura:cobertura -Dcobertura.report.f
15                     },
16                     sonar: {
17                         sh "mvn sonar:sonar -Dsonar.host.url=${env.SONARQUBE_HOST}"
18                     }
19                 )
20             }
21             post {
22                 always {
23                     junit '**/target/*-reports/TEST-*.xml'
24                     step([$class: 'CoberturaPublisher', coberturaReportFile: 'target/site,
25                 }
26             }
27         }
28         stage('deploy') {
29             steps {
30                 sh "mvn deploy -DskipTests"
31             }
32         }
33     }
34 }

```

This is what going to happen:

```

1  [INFO] BUILD FAILURE
2  [INFO] -----
3  [INFO] Total time: 9.265 s
4  [INFO] Finished at: 2017-07-16T18:39:51Z
5  [INFO] Final Memory: 17M/95M
6  [INFO] -----
7  [ERROR] Failed to execute goal org.apache.maven.plugins:maven-deploy-plugin:2.7:deploy (de
8  [ERROR]
9  [ERROR] To see the full stack trace of the errors, re-run Maven with the -e switch.
10 [ERROR] Re-run Maven using the -X switch to enable full debug logging.
11 [ERROR]
12 [ERROR] For more information about the errors and possible solutions, please read the foll
13 [ERROR] [Help 1] http://cwiki.apache.org/confluence/display/MAVEN/MojoExecutionException
14 [Pipeline] }

```

The job fails, but Maven is clever enough to figure out the reason, and it prompted us with a message that the `distributionManagement` section in the pom file is missing. So the Maven deploy plugin doesn't basically know where to deploy the artifact, would you know?

Chapter 2: Configuring Maven pom File

Let's configure the missing part in the pom file inside the project. If you are not using your own git project for this tutorial, you won't be able to amend the POM file, as it is in my repository and you don't have access, unfortunately. But you can switch to another project in that case, which I prepared for you. So please create another pipeline job for the project `maze-explorer`, which has the necessary pom changes. If you are using your own Git project, just ignore this note. Obviously, I could have used just a branch in the same project, but let's have a couple jobs in Jenkins!


Creating a new pipeline:

:8080/view/all/newJob


Enter an item name

maze-explorer


» Required field


Freestyle project

This is the central feature of Jenkins. Jenkins will build used for something other than software build.

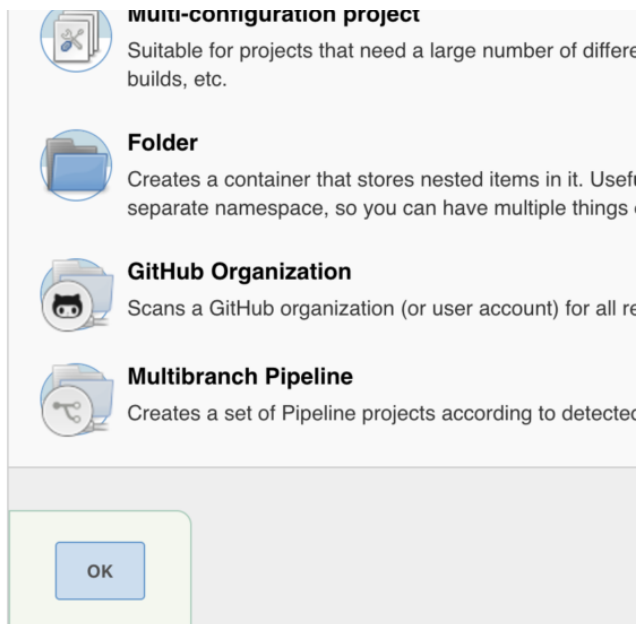

Pipeline

Orchestrates long-running activities that can span multi and/or organizing complex activities that do not easily f

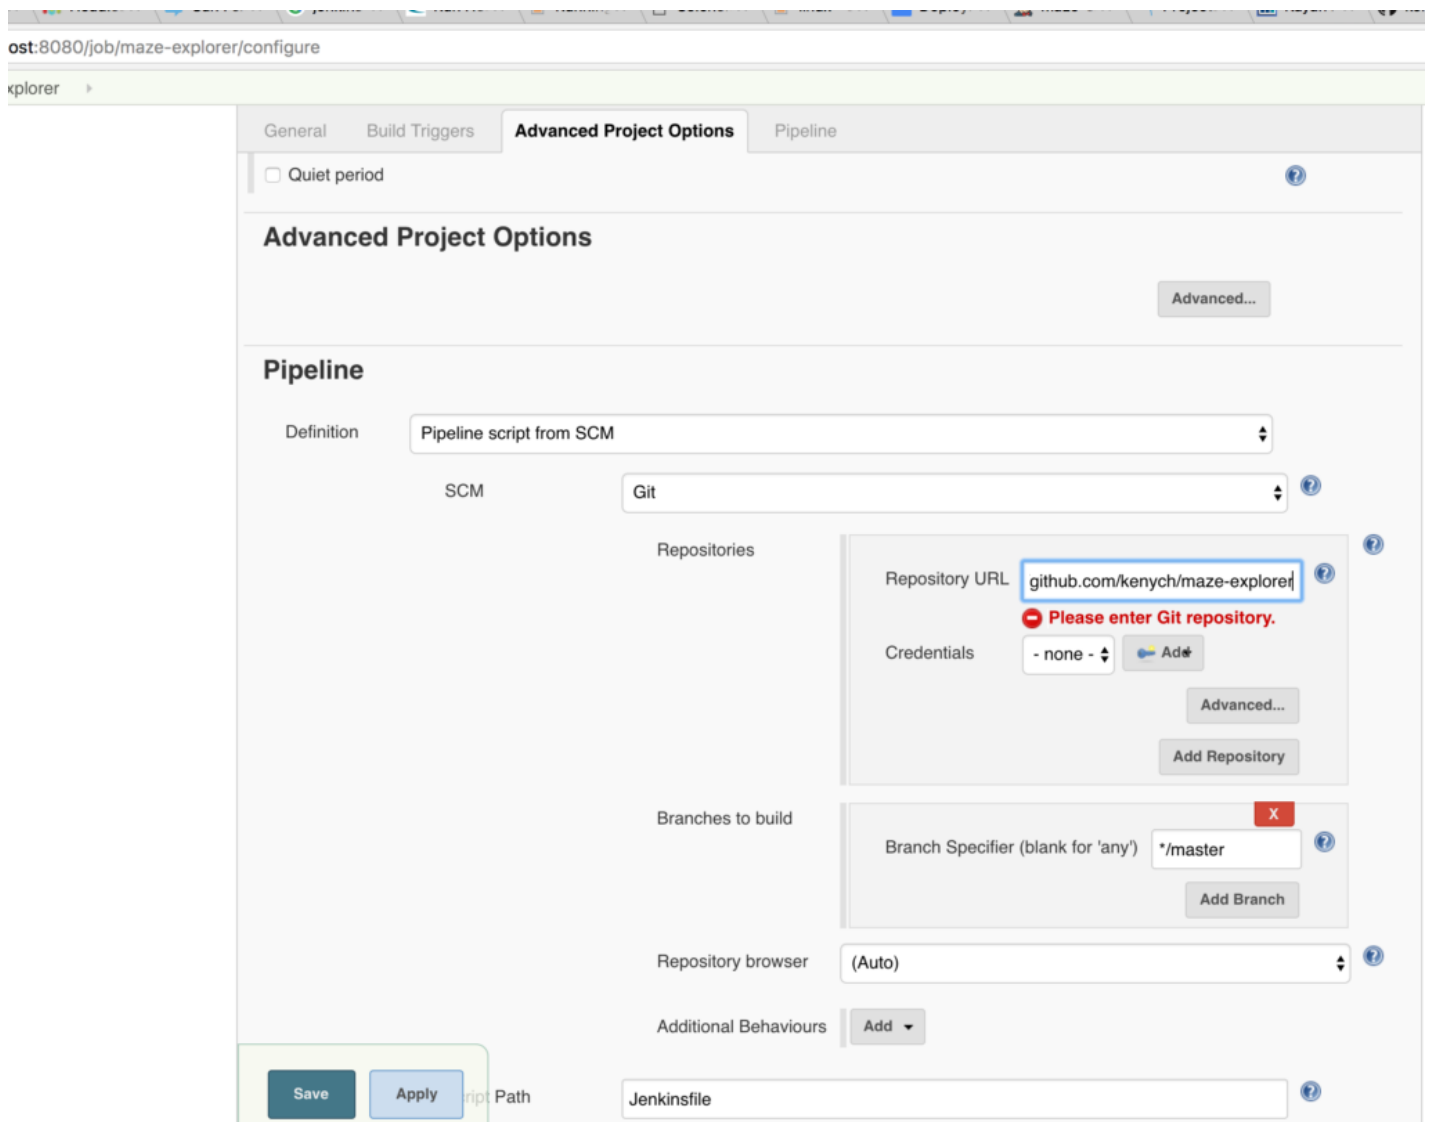

External Job

This type of job allows you to record the execution of a that you can use Jenkins as a dashboard of your existi

Multi-configuration project



Configuring the repository:



Here is how your changes in the POM will look, We added distributionManagement to the project section:

```
<project xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance" xmlns="http://maven.apache.org/POM/4.0.0"
  <modelVersion>4.0.0</modelVersion>
```



```

2  .....
3
4  <groupId>kayan</groupId>
5  <artifactId>maze</artifactId>
6  <version>1.0-SNAPSHOT</version>
7  <packaging>jar</packaging>
8
9  <url>http://maven.apache.org</url>
10
11 <properties>
12     <project.build.sourceEncoding>UTF-8</project.build.sourceEncoding>
13 </properties>
14
15 <distributionManagement>
16     <snapshotRepository>
17         <id>artifactory</id>
18         <name>artifactory</name>
19         <url>${artifactory_url}</url>
20     </snapshotRepository>
21 </distributionManagement>
22
23
24 (the rest of POM file)

```

The URL should point to the URL of Artifactory, so we need to pass it to Maven through Jenkins as an environment variable, just like we did for Sonar. Obviously it can't be static as the IP of the host can change.

If you run the job now, it is still going to fail, even if it has distributionManagement configured:

```

1  [INFO] BUILD FAILURE
2  [INFO] -----
3  [INFO] Total time: 4.459 s
4  [INFO] Finished at: 2017-07-16T20:13:34Z
5  [INFO] Final Memory: 13M/137M
6  [INFO] -----
7  [ERROR] Failed to execute goal org.apache.maven.plugins:maven-deploy-plugin:2.7:deploy (de
8  [ERROR]
9  [ERROR] To see the full stack trace of the errors, re-run Maven with the -e switch.
10 [ERROR] Re-run Maven using the -X switch to enable full debug logging.

```





This is because we need to have credentials in order to deploy. The default credentials of JFrog are “admin:password.” Try to log in and check. In order to pass credentials to the deployment plugin, we need to set them in the Maven settings.xml file.

Chapter 3: Configuring the Maven Settings File

If we have Maven installed locally, then we can run it first just to check that the deployment actually works with our configuration, and then we can start looking at how to configure it with Jenkins. Let's check what we have in the settings file:

```
1 mvn help:effective-settings
```

You will see something similar if the settings file is absent or empty.

```
1 <settings xmlns="http://maven.apache.org/POM/4.0.0" xmlns:xsi="http://www.w3.org/2001/XMLSchema-
2 <localRepository xmlns="http://maven.apache.org/SETTINGS/1.1.0"/>/Users/kayanazimov/.m2/r
3 <pluginGroups xmlns="http://maven.apache.org/SETTINGS/1.1.0">
4   <pluginGroup>org.apache.maven.plugins</pluginGroup>
5   <pluginGroup>org.codehaus.mojo</pluginGroup>
6 </pluginGroups>
7 </settings>
```

Now let's go to /Users/YOUR_USER_NAME_HERE/.m2/ and change or create settings.xml as below. In order to pass credentials of the repo, we need to add a server in the servers section of the Maven settings:

```
1 <?xml version="1.0" encoding="UTF-8"?>
2 <settings xsi:schemaLocation="http://maven.apache.org/SETTINGS/1.1.0 http://maven.apache.org/
3 <servers>
4   <server>
5     <id>artifactory</id>
6     <username>admin</username>
7     <password>password</password>
8   </server>
9 </servers>
10 </settings>
```

Please note, the id should be same as the id you used earlier in the pom file for snapshotRepository.

You can check if Maven is picking it up:

```
1 mvn help:effective-settings
```

Now clone the project, or if you're using your own project, just run it from the project folder (use the Artifactory port, which will be shown when running the runall script):

```

1  mvn deploy -DskipTests -Dartifactory_url=http://localhost:8082/artifactory/example-repo-local
2
3  [INFO] Scanning for projects...
4  [INFO]
5  [INFO] -----
6  [INFO] Building maze 1.0-SNAPSHOT
7  [INFO] -----
8  [INFO]
9  [INFO] --- maven-resources-plugin:2.6:resources (default-resources) @ maze ---
10 [INFO] Using 'UTF-8' encoding to copy filtered resources.
11 [INFO] skip non existing resourceDirectory /Users/kayanazimov/workspace/learn/maze-explorer
12 [INFO]
13 [INFO] --- maven-compiler-plugin:2.5.1:compile (default-compile) @ maze ---
14 [INFO] Nothing to compile - all classes are up to date
15 [INFO]
16 [INFO] --- maven-resources-plugin:2.6:testResources (default-testResources) @ maze ---
17 [INFO] Using 'UTF-8' encoding to copy filtered resources.
18 [INFO] Copying 10 resources
19 [INFO]
20 [INFO] --- maven-compiler-plugin:2.5.1:testCompile (default-testCompile) @ maze ---
21 [INFO] Nothing to compile - all classes are up to date
22 [INFO]
23 [INFO] --- maven-surefire-plugin:2.12.4:test (default-test) @ maze ---
24 [INFO] Tests are skipped.
25 [INFO]
26 [INFO] --- maven-jar-plugin:2.4:jar (default-jar) @ maze ---
27 [INFO] Building jar: /Users/kayanazimov/workspace/learn/maze-explorer/target/maze-1.0-SNAPSHOT.jar
28 [INFO]
29 [INFO] --- maven-install-plugin:2.4:install (default-install) @ maze ---
30 [INFO] Installing /Users/kayanazimov/workspace/learn/maze-explorer/target/maze-1.0-SNAPSHOT.jar to /Users/kayanazimov/.m2/repository/kayan/maze/1.0-SNAPSHOT/maze-1.0-SNAPSHOT.jar
31 [INFO] Installing /Users/kayanazimov/workspace/learn/maze-explorer/pom.xml to /Users/kayanazimov/.m2/repository/kayan/maze/1.0-SNAPSHOT/maze-1.0-SNAPSHOT.pom
32 [INFO]
33 [INFO] --- maven-deploy-plugin:2.7:deploy (default-deploy) @ maze ---
34 Downloading: http://localhost:8082/artifactory/example-repo-local/kayan/maze/1.0-SNAPSHOT/maze-1.0-SNAPSHOT.jar
35 Uploading: http://localhost:8082/artifactory/example-repo-local/kayan/maze/1.0-SNAPSHOT/maze-1.0-SNAPSHOT.jar
36 Uploaded: http://localhost:8082/artifactory/example-repo-local/kayan/maze/1.0-SNAPSHOT/maze-1.0-SNAPSHOT.jar
37 Uploading: http://localhost:8082/artifactory/example-repo-local/kayan/maze/1.0-SNAPSHOT/maze-1.0-SNAPSHOT.pom

```

```

38  Uploaded: http://localhost:8082/artifactory/example-repo-local/kayan/maze/1.0-SNAPSHOT/maz
39  Downloading: http://localhost:8082/artifactory/example-repo-local/kayan/maze/maven-metadat
40  Uploading: http://localhost:8082/artifactory/example-repo-local/kayan/maze/1.0-SNAPSHOT/maz
41  Uploaded: http://localhost:8082/artifactory/example-repo-local/kayan/maze/1.0-SNAPSHOT/mav
42  Uploading: http://localhost:8082/artifactory/example-repo-local/kayan/maze/maven-metadata.
43  Uploaded: http://localhost:8082/artifactory/example-repo-local/kayan/maze/maven-metadata.
44  [INFO] -----
45  [INFO] BUILD SUCCESS
46  [INFO] -----
47  [INFO] Total time: 1.828 s
48  [INFO] Finished at: 2017-07-18T09:39:52+01:00
49  [INFO] Final Memory: 13M/207M
50  [INFO] -----
51
52  → maze-explorer git:(master)

```

Now check the Artifactory:

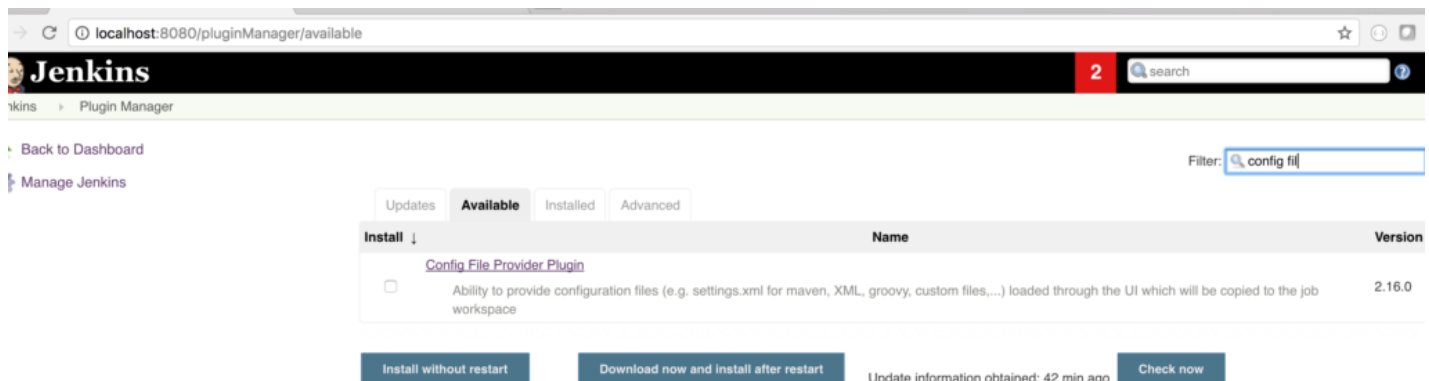
The screenshot displays the JFrog Artifactory web interface. The browser is set to 'example-repo-local'. The left sidebar shows a tree view of the repository structure, including 'kayan/maze/1.0-SNAPSHOT' and 'maze-1.0-20170716.203402-1.jar'. The right panel shows the 'General' tab with repository details like Name, Package Type, Repository Path, and Created date.

As you can see we successfully deployed it, Yay! Now it is time to prepare this step in Jenkins.

Chapter 4: Using the Config File Provider Plugin

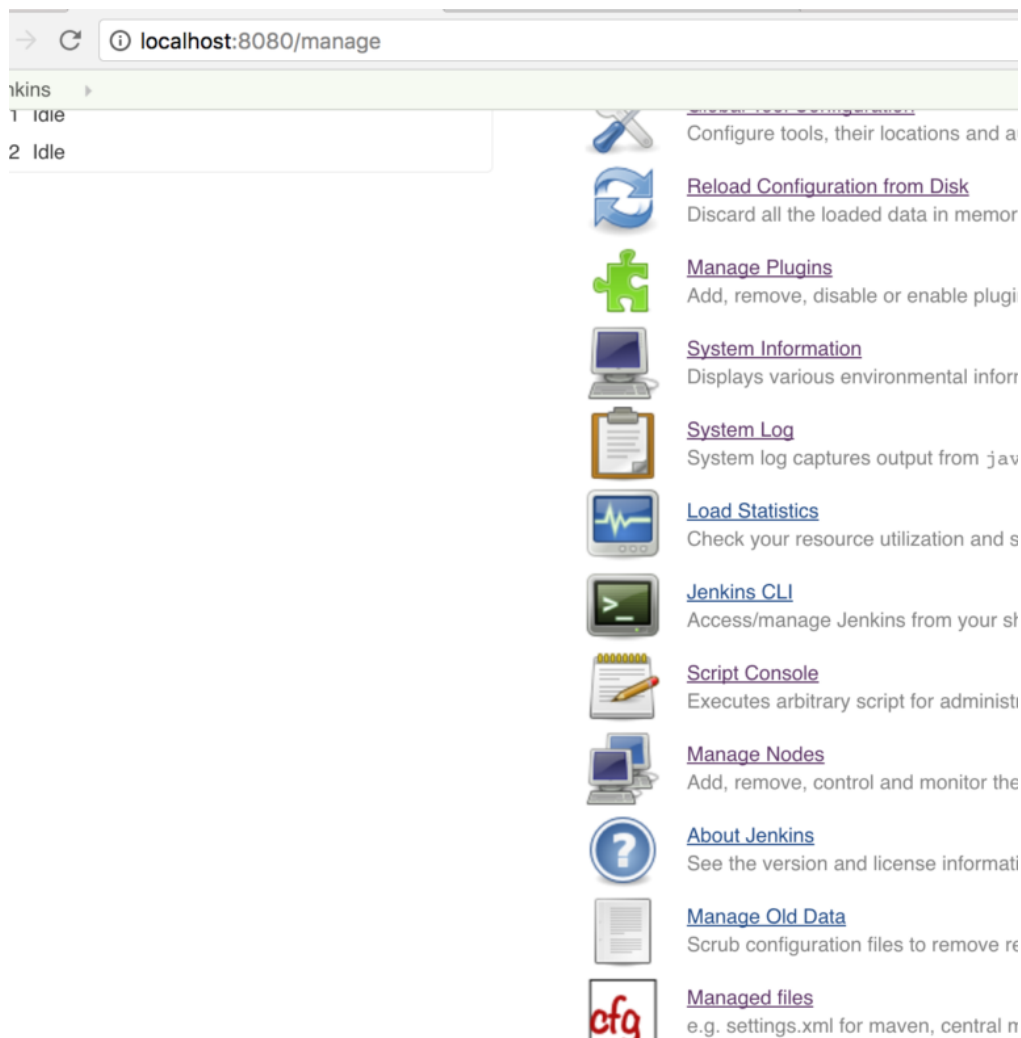
for Persistence of Maven Settings

In order to apply the settings file to Maven in Jenkins, we need the config file provider plugin, which lets us preserve multiple settings files (please be aware, we might need more than one settings depending on a project running the job in real life). Now let's install the plugin manually first:

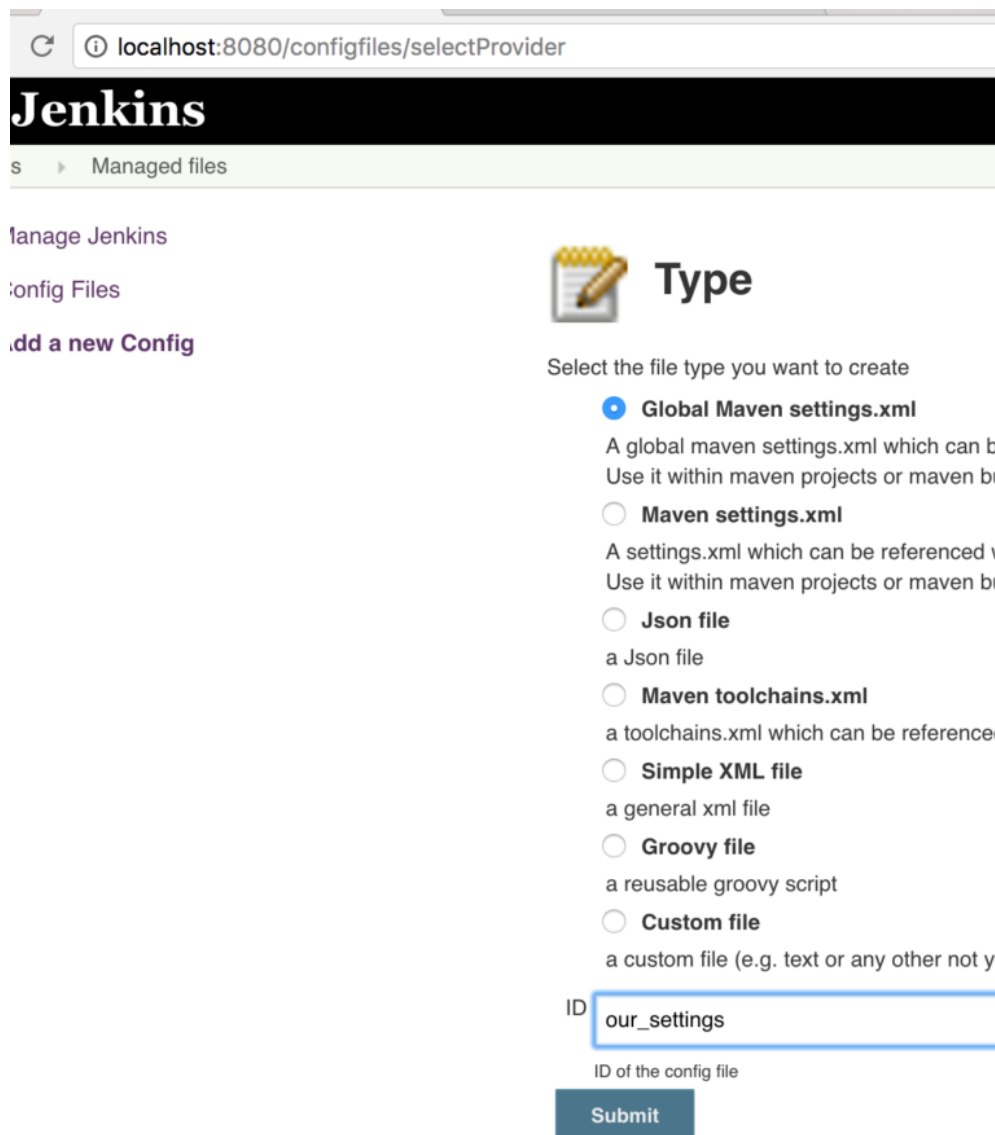


We need to set up the Maven settings file in the plugin. Go to manage Jenkins, managed files, and add new config, set the id to "our_settings" and copy the content from the settings.xml we used before:

Click Managed files:



Select Global Maven Settings and set id:



The screenshot shows the Jenkins web interface at the URL `localhost:8080/configfiles/selectProvider`. The page title is "Jenkins" and the breadcrumb is "Managed files". On the left sidebar, there are links for "Manage Jenkins", "Config Files", and "Add a new Config". The main content area is titled "Type" with a notepad icon. It asks the user to "Select the file type you want to create". There are seven radio button options: "Global Maven settings.xml" (selected), "Maven settings.xml", "Json file", "Maven toolchains.xml", "Simple XML file", "Groovy file", and "Custom file". Each option has a brief description. Below the options is an "ID" input field containing the text "our_settings". A "Submit" button is at the bottom right.

localhost:8080/configfiles/selectProvider

Jenkins

Managed files

Manage Jenkins

Config Files

Add a new Config

Type

Select the file type you want to create

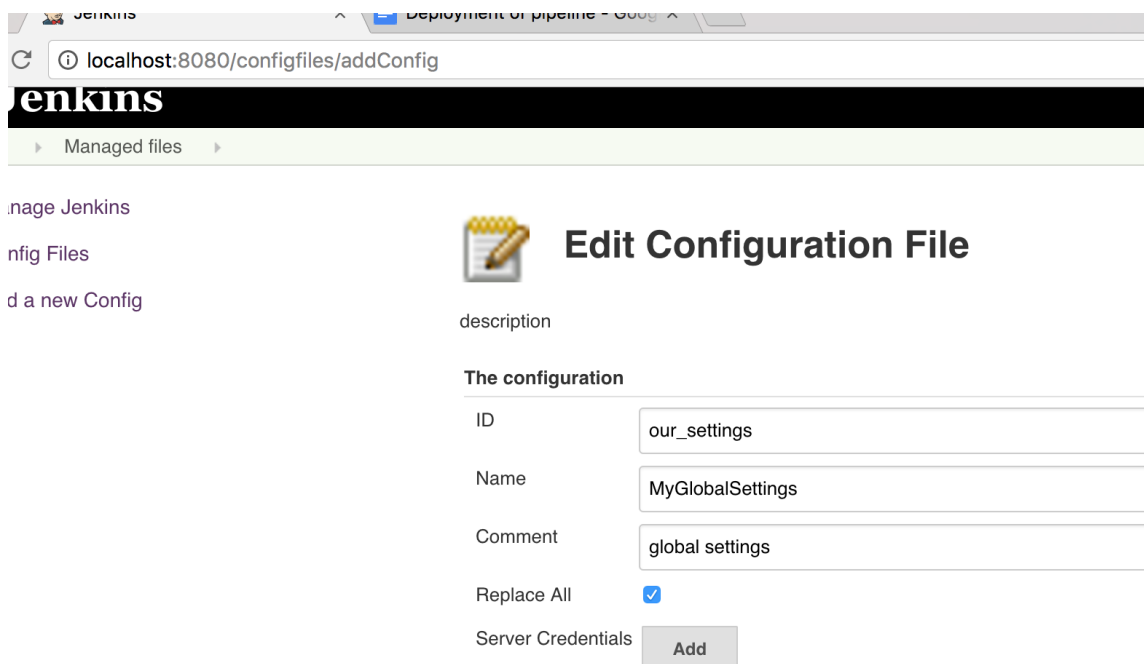
- ☒ **Global Maven settings.xml**
A global maven settings.xml which can be used within maven projects or maven build scripts
- ☐ **Maven settings.xml**
A settings.xml which can be referenced within maven projects or maven build scripts
- ☐ **Json file**
A Json file
- ☐ **Maven toolchains.xml**
A toolchains.xml which can be referenced within maven projects or maven build scripts
- ☐ **Simple XML file**
A general xml file
- ☐ **Groovy file**
A reusable groovy script
- ☐ **Custom file**
A custom file (e.g. text or any other not yet supported file type)

ID

ID of the config file

Submit

And set content:



The screenshot shows the Jenkins web interface at the URL `localhost:8080/configfiles/addConfig`. The page title is "Jenkins" and the breadcrumb is "Managed files". On the left sidebar, there are links for "Manage Jenkins", "Config Files", and "Add a new Config". The main content area is titled "Edit Configuration File" with a notepad icon. It has a "description" label. Below that is a section titled "The configuration" with four input fields: "ID" (containing "our_settings"), "Name" (containing "MyGlobalSettings"), "Comment" (containing "global settings"), and "Replace All" (checked). There is also a "Server Credentials" label and an "Add" button.

localhost:8080/configfiles/addConfig

Jenkins

Managed files

Manage Jenkins

Config Files

Add a new Config

Edit Configuration File

description

The configuration

ID

Name

Comment

Replace All ☒

Server Credentials

Content

```

1 <?xml version="1.0" encoding="UTF-8"?>
2 <settings xsi:schemaLocation="http://maven.apache.org/xsd/maven-4.0.0.xsd" xmlns="http://maven.apache.org/2002/CI
3   xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance">
4   <servers>
5     <server>
6       <id>artifactory</id>
7       <username>admin</username>
8       <password>password</password>
9     </server>
10  </servers>
11 </settings>
12

```

Now we can use the Config File Provider plugin in our pipeline. Please Replay the job and update the pipeline as below:

```

1 pipeline {
2   agent any
3
4   tools {
5     jdk 'jdk8'
6     maven 'maven3'
7   }
8
9   stages {
10    stage('install and sonar parallel') {
11      steps {
12        parallel(
13          install: {
14            sh "mvn -U clean test cobertura:cobertura -Dcobertura.report.file=target/site/cobertura.xml"
15          },
16          sonar: {
17            sh "mvn sonar:sonar -Dsonar.host.url=${env.SONARQUBE_HOST}"
18          }
19        )
20      }
21      post {
22        always {
23          junit '**/target/*-reports/TEST-*.xml'
24          step([$class: 'CoberturaPublisher', coberturaReportFile: 'target/site/cobertura.xml'])
25        }
26      }
27    }
28    stage('deploy'){
29      steps{
30        configFileProvider([configFile(fileId: 'our_settings', variable: 'SETTINGS')])
31        sh "mvn -s $SETTINGS deploy -DskipTests -Dartifactory.url=${env.ARTIFACTORY_URL}"
32      }
33    }
34  }
35 }

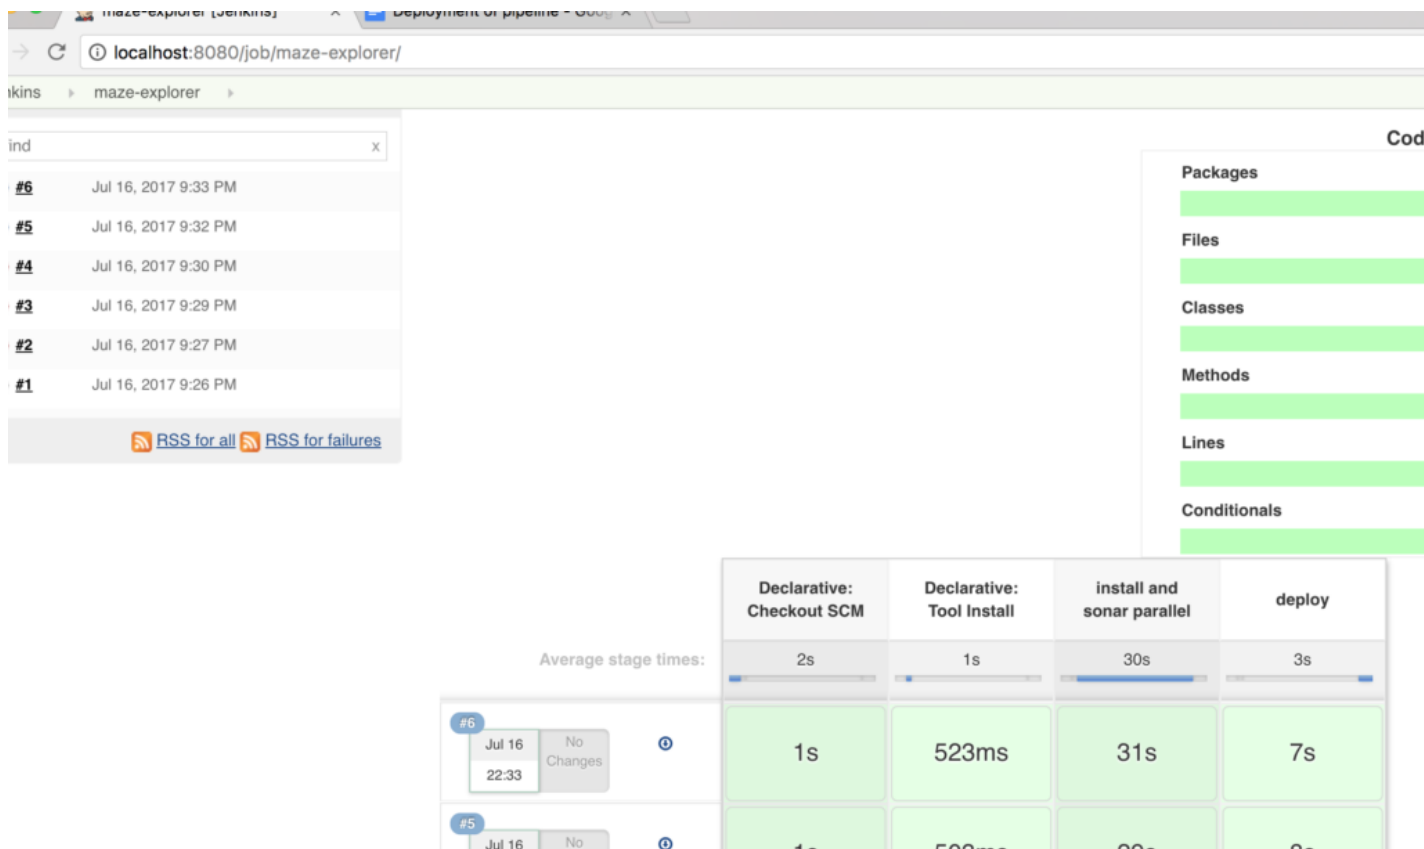
```

```

31      }
32    }
33  }
34  }
35  }
36  }

```

Let's run the build:



If you are lucky, you will get the screen above. Otherwise, read what can go wrong with containers when running out of memory.

Chapter 5: Dockerizing the Installation and Configuration Process

Now let's dockerize the plugin installation and configuration. Create `mvn_settings.groovy` file, copy it to the groovy folder that we created in the first part of the tutorial, and set content as below:

```

1  import jenkins.model.*
2  import org.jenkinsci.plugins.configfiles.maven.*
3  import org.jenkinsci.plugins.configfiles.maven.security.*
4
5  def store = Jenkins.instance.getExtensionList('org.jenkinsci.plugins.configfiles.GlobalCor
6
7
8  println("Setting maven settings xml")
9
10

```

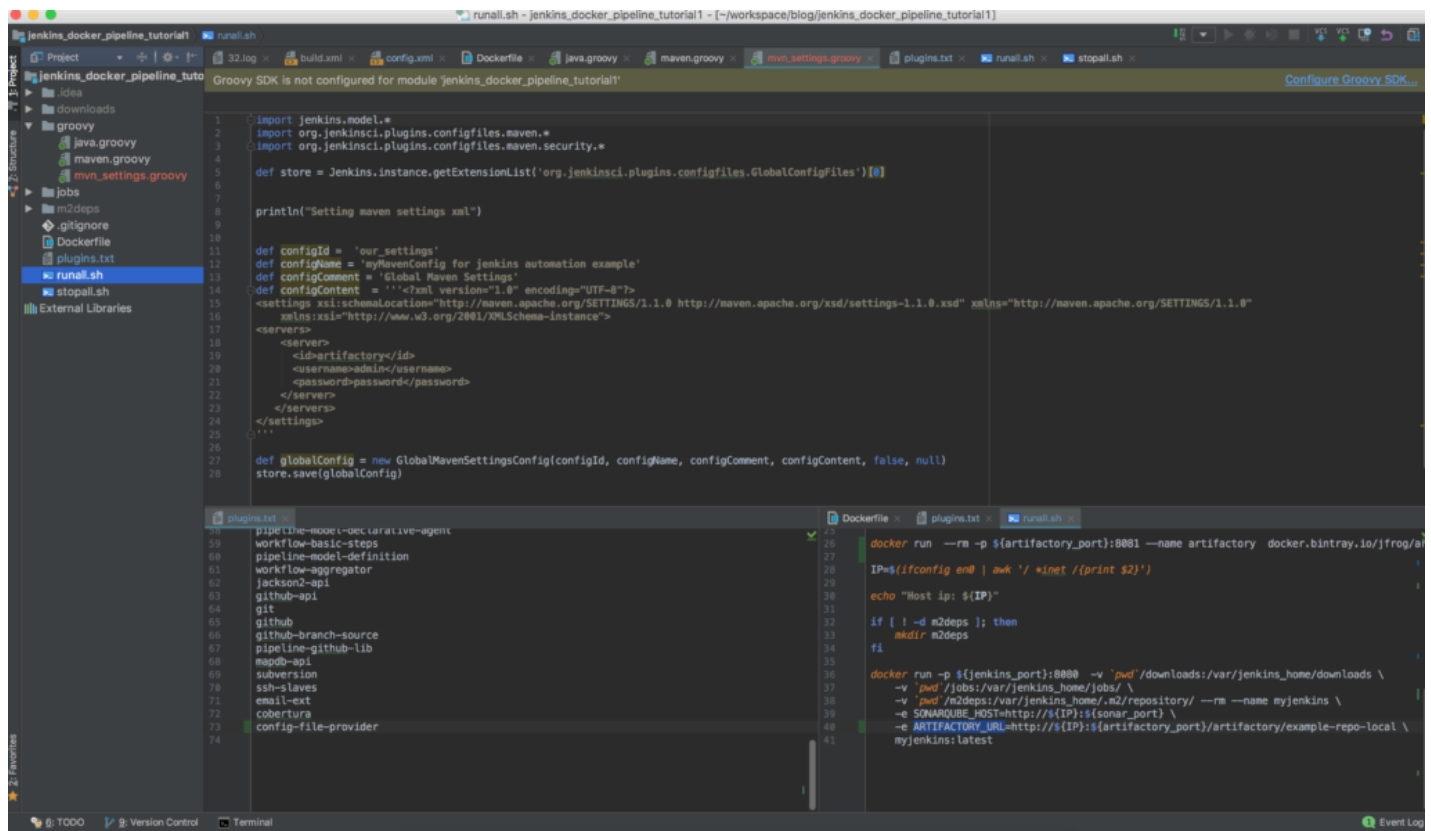


```


11 def configId = 'our_settings'
12 def configName = 'myMavenConfig for jenkins automation example'
13 def configComment = 'Global Maven Settings'
14 def configContent = '''<?xml version="1.0" encoding="UTF-8"?>
<settings xsi:schemaLocation="http://maven.apache.org/SETTINGS/1.1.0 http://maven.apache.c
15 <
16 <servers>
17     <server>
18         <id>artifactory</id>
19         <username>admin</username>
20         <password>password</password>
21     </server>
22 </servers>
23 </settings>
24 '''
25
26 def globalConfig = new GlobalMavenSettingsConfig(configId, configName, configComment, confi
27 store.save(globalConfig)

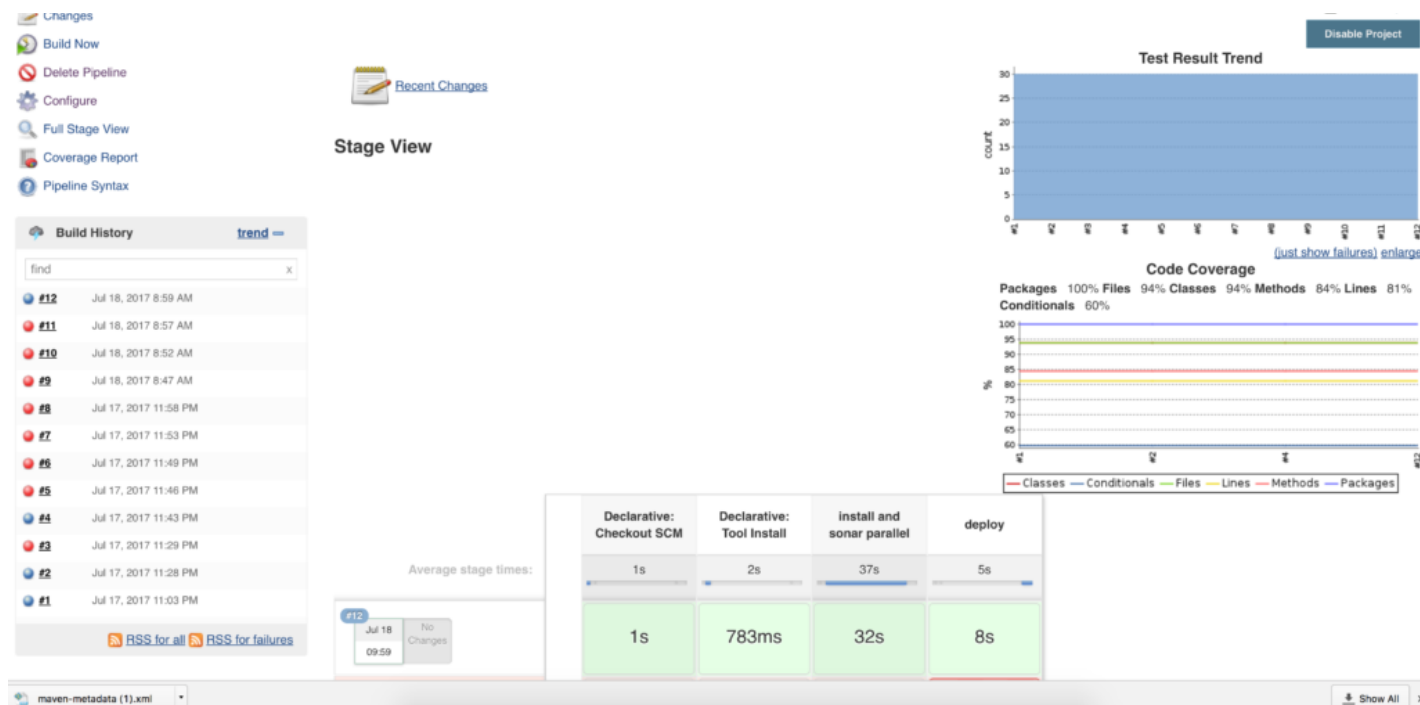
```

And don't forget to install the “config-file-provider” plugin, just add it to the plugins.txt like in this screenshot:



Time to rebuild and rerun your Jenkins, just run the runall script. Once Jenkins is ready, run the build:





We now have the CI pipeline for our project and almost fully automated Jenkins. Why almost?



Because we need to hide the passwords firstly. But I am running late, and my girlfriend has already started complaining...



But I promise we will look at how to use encrypted passwords in Jenkins very soon in the next session.

Again, If you were lazy and don't like "hands on" stuff, just clone the complete code for the second part of the tutorial, run it, enjoy it, and share it if you like:

```
git clone https://github.com/kenych/dockerizing-jenkins-part-2 && cd dockerizing-jenkins-1
```

Hope you managed to run everything without any issues and enjoyed this tutorial. Should you have any issue don't hesitate to comment and I will try to help.

Download the 'Practical Blueprint to Continuous Delivery' to learn how Automatic Release Automation can help you begin or continue your company's digital transformation.

Like This Article? Read More From DZone



Unleash the DevOps!



Setting Up Jenkins to Deploy to Heroku



Orchestrating Workflows With Jenkins and Docker



**Free DZone Refcard
Getting Started With Docker**

Topics: DOCKER , DEPLOY , ARTIFACTORY , JENKINS

Opinions expressed by DZone contributors are their own.

Get the best of DevOps in your inbox.

Stay updated with DZone's bi-weekly DevOps Newsletter. [SEE AN EXAMPLE](#)

[SUBSCRIBE](#)

DevOps Partner Resources

The Ultimate Monitoring Tools Landscape

Papertrail



Setting Up the Critical DevOps Role of Enterprise Release Management

Plutora



The DevOps Journey - From Waterfall to Continuous Delivery

Sauce Labs



Redefine application release processes and unify your DevOps strategy.

Automatic



DDD, Part 2: DDD Building Blocks

by M Yauri Maulana at-Tamimi · Dec 30, 17 · DevOps Zone

Learn more about how CareerBuilder was able to resolve customer issues 5x faster by using Scalyr, the fastest log management tool on the market.

It's been a while since I wrote the first part of my DDD series. In this second part, I will continue with one of the most important things to know about DDD before you go further with your implementation. First things first: you must know the DDD building blocks below:

1. Entities
2. Value objects
3. Aggregate roots
4. Repositories
5. Factories
6. Services

Fasten your seatbelts!. We're going through the details of those blocks now.

1. Entities

An entity is a plain object that has an identity (ID) and is potentially mutable. Each entity is uniquely identified by an ID rather than by an attribute; therefore, two entities can be considered equal (identifier equality) if both of them have the same ID even though they have different attributes. This means that the state of the entity can be changed anytime, but as long as two entities have the same ID, both are considered equal regardless what attributes they have.

2. Value Objects

Value objects are immutable. They have no identity (ID) like we found in entity. Two value objects can be considered equal if both of them have the same type and the same attributes (applied to all of its attributes).

There are often uses for a thing like message passing and in fact, this is particularly useful in the API layer within an onion architecture to expose your domain concepts without necessarily exposing the immutable aspect.

Some benefits of value objects:

1. The compound of value objects can swallow lots of computational complexity.
2. Entities can be released from logic complexity.
3. Improve extensibility, especially for testability and concurrency issues if using correctly.

3. Aggregate Roots

Aggregate root is an entity that binds together with other entities. Moreover, aggregate root is actually a part of aggregate (collection/cluster of associated objects that are treated as a single unit for the purpose of data changes). Thus, each aggregate actually consists of an aggregate root and a boundary. For example, the relationship between Order and OrderLineItem within SalesOrderDomain can be considered as an aggregate where Order acts as the aggregate root, while the OrderLineItem is the child of Order within SalesOrder boundary.

One of the key features of an aggregate root is that the external objects are not allowed to hold a reference to an aggregate root child entities. Thus, if you need access to one of the aggregate root child entities (AKA aggregate), then you must go through the aggregate root (i.e. you can't access the child directly).

The other thing is that all operations within the domain should, where possible, go through an aggregate root. Factories, repositories, and services are some exceptions to this, but whenever possible, if you can create or require that an operation goes through the aggregate root, that's going to be better.

4. Repositories

Repositories are mostly used to deal with the storage. They are actually one of the most important concepts on the DDD because they have abstracted away a lot of the storage concerns (i.e. some form/mechanism of storage).

The repository implementation could be a file-based storage, or database (SQL-/NoSQL-based), or any other thing that is related to storage mechanism, such as caching. Any combination of those is also possible.

A repository should not be confused with the data store. A repository job is to store aggregate roots. Underneath that, repositories implementation may actually have to talk to multiple different storage locations in order to construct the aggregates. Thus, a single aggregate root might be drawn from a REST API, as well as a database or files. You may wrap those in something called the data store, but the repository is sort of a further layer of an abstraction on top of all those individual data stores. Usually, I implement the repository as an interface within the domain/domain services layer within onion architecture, and then the implementation logic of the repository interface is going to be defined in the infrastructure layer.

5. Factories

The factories are used to give an abstraction to the object construction (see factory design pattern from GOF).

A factory can also potentially return an aggregate root or an entity, or perhaps a value object. Often times, when you need a factory method for an aggregate root, it will be rolled into the repository. So, your repository might have a finder create method on it.

Usually, factories also implemented as an interface within the domain/domain services layer with the implementation logic will be defined in the infrastructure layer.

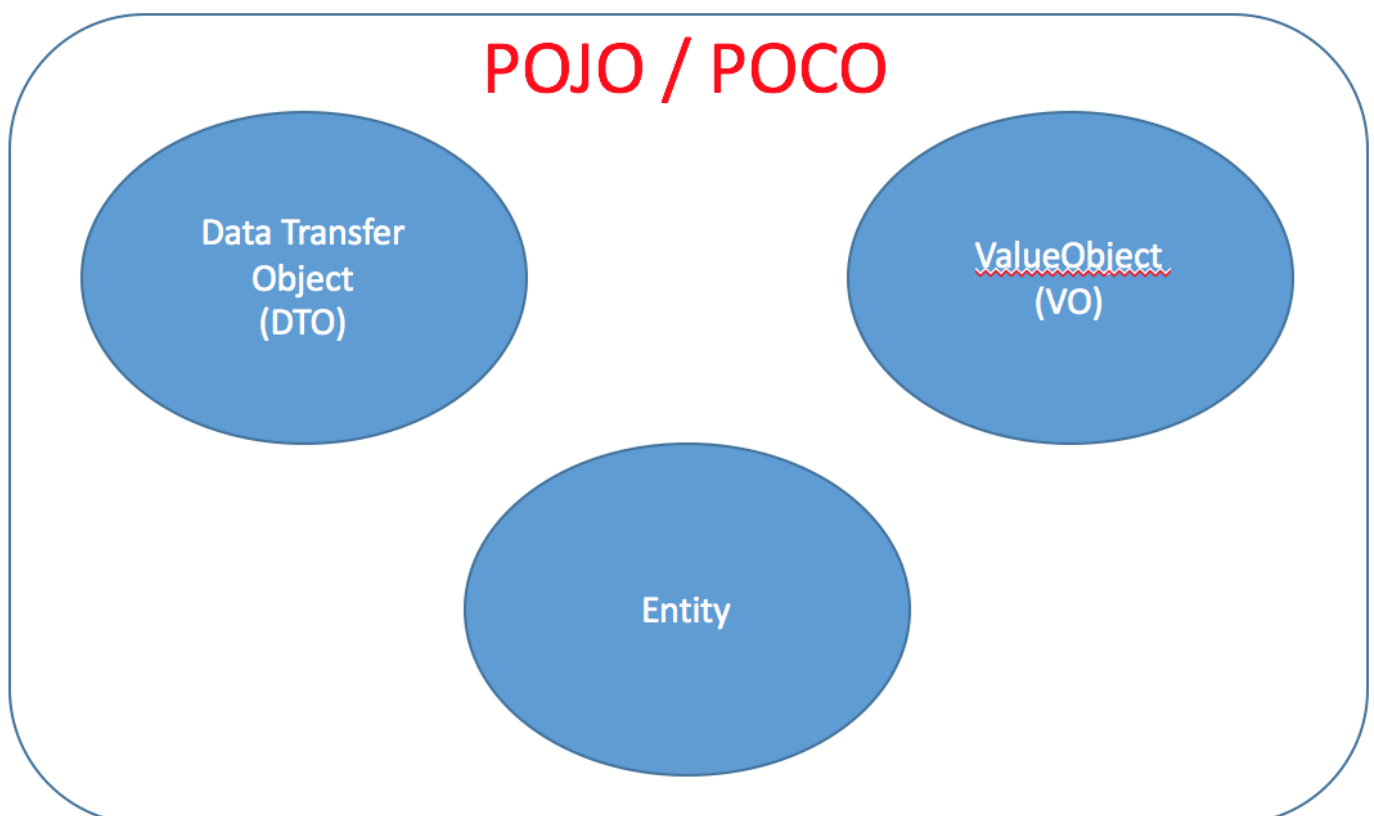
6. Services

A service basically exists to provide a home for operations that don't quite fit into an aggregate root. As an example, when you have an operation and don't know which aggregate root it goes into, perhaps it operates on multiple aggregate roots or maybe it doesn't belong to any existing aggregate root. Then, you can put the logic into a service. However, don't be rushed to put everything into a service. First and foremost, it's better to carefully analyze whether the operation fits into one of the existing aggregate roots. If you couldn't find the aggregate root, it's subsequently better to ask yourself if you have missed one aggregate root, or perhaps there are domain concepts that you haven't considered that should be brought into your domain before you put the operation into a service.

Other Important Things

I often found that many developers use term VO (value objects) and DTO (data transfer object) interchangeably. They think both are just the same. This is quite annoying for me. I'd like to clarify here that both refer to the different things.

As depicted in the picture below, VO and DTO are subsets of a POJO/POCO. An entity is also a subset of POJO/POCO.



In the above depiction, POJO and POCO can be interchangeably used. Both are referring to similar things. Both are just domain objects that mostly represent the domain/business object within the business application.

The term POJO (plain old Java object) was coined by Martin Fowler and very popular in Java community, while POCO (plain old CLR object/plain old class object) is widely used in the dotNet.

As mentioned earlier, DTO, VO, and entity are just a subset of POJO/POCO. However, they are really different things as described below:

	Has Data	Has Logic	Has ID
Data Transfer Object	Yes	No	No
Value Object	Yes	Yes	No
Entity	Yes	Yes	Yes

DTO is merely a stupid data container (only holds data without any logic). Thus, it's anemic in general (only contains attributes and getter/setter). DTO is absolutely immutable. Usually, we use DTO to transfer the object between layers and tiers in one single application or between application to application or JVM to JVM (mostly useful between networks to reduce multiple network call).

VO is also immutable, but what makes it different than DTO is that VO also contains logic.

That's all for now. Read the next part here.

Find out more about how Scalyr built a proprietary database that does not use text indexing for their log management tool.

Like This Article? Read More From DZone



Value Objects



Architect Your iOS App for Easy Backend Replacement - Part I: Functional Domain Design




Learning Angular 2: Tour of Heroes Tutorial, HTTP



Free DZone Refcard Getting Started With Docker

Topics: ENTITIES, DEVOPS, REPOSITORIES, FACTORIES, SERVICES, VALUE OBJECTS, AGGREGATE ROOTS, TUTORIAL

Published at DZone with permission of M Yauri Maulana at-Tamimi. [See the original article here.](#) 
Opinions expressed by DZone contributors are their own.
